

Peak Allocation Events

- A Model Parameters
- B Optimizer States (AdamW)
- C cuBLAS Workspace
- D Pipeline Boundary Copies
- E Gradient Boundary Copies
- F Parameter Gradients
- G Recomputed Activations
- H Final Backward Tensors
- I OTHERS

Idx	Cat	Size(B)	File	Line	Code
1	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
2	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
3	A	7077888	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
4	A	9216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
5	A	2359296	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
6	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
7	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
8	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
9	A	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
10	A	12288	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
11	A	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
12	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
13	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
14	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
15	A	7077888	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
16	A	9216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
17	A	2359296	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
18	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
19	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
20	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
21	A	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
22	A	12288	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
23	A	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
24	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
25	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
26	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
27	A	7077888	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
28	A	9216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
29	A	2359296	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
30	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
31	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
32	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
33	A	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
34	A	12288	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
35	A	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
36	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
37	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
38	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
39	A	7077888	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
40	A	9216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
41	A	2359296	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
42	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
43	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
44	A	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(
45	A	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	pipe_model = GPipe(

[illegible]

Idx	Cat	Size(B)	File	Line	Code
104	B	9216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
105	B	9216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
106	B	2359296	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
107	B	2359296	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
108	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
109	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
110	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
111	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
112	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
113	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
114	B	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
115	B	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
116	B	12288	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
117	B	12288	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
118	B	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
119	B	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
120	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
121	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
122	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
123	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
124	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
125	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
126	B	7077888	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
127	B	7077888	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
128	B	9216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
129	B	9216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
130	B	2359296	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
131	B	2359296	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
132	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
133	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
134	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
135	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
136	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
137	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
138	B	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
139	B	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
140	B	12288	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
141	B	12288	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
142	B	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
143	B	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
144	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
145	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	optimizer.step()
146	D	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	236	logits = pipe_model(input_ids)
147	G	8519680	~/mew_mem_1/gpt2_gpipe_memory_profile.py	90	qkv = self.c_attn(x)
148	D	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	236	logits = pipe_model(input_ids)
149	D	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	236	logits = pipe_model(input_ids)
150	H	9437184	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260-packages/torch/pa	260	chgnpe.attn.backward(tensors, grad_output)
151	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260-packages/torch/pa	260	chgnpe.attn.backward(tensors, grad_output)
152	H	9437184	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260-packages/torch/pa	260	chgnpe.attn.backward(tensors, grad_output)
153	H	12288	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260-packages/torch/pa	260	chgnpe.attn.backward(tensors, grad_output)
154	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260-packages/torch/pa	260	chgnpe.attn.backward(tensors, grad_output)
155	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260-packages/torch/pa	260	chgnpe.attn.backward(tensors, grad_output)

Idx	Cat	Size(B)	File	Line	Code
185	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	self.packages/torch_gpipe_attn_grad.backward(tensors, grad_output)
186	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	self.packages/torch_gpipe_attn_grad.backward(tensors, grad_output)
187	H	9437184	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	self.packages/torch_gpipe_attn_grad.backward(tensors, grad_output)
188	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	self.packages/torch_gpipe_attn_grad.backward(tensors, grad_output)
189	H	9437184	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	self.packages/torch_gpipe_attn_grad.backward(tensors, grad_output)
190	H	12288	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	self.packages/torch_gpipe_attn_grad.backward(tensors, grad_output)
191	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	self.packages/torch_gpipe_attn_grad.backward(tensors, grad_output)
192	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	self.packages/torch_gpipe_attn_grad.backward(tensors, grad_output)
193	H	2359296	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	self.packages/torch_gpipe_attn_grad.backward(tensors, grad_output)
194	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	self.packages/torch_gpipe_attn_grad.backward(tensors, grad_output)
195	H	7077888	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	self.packages/torch_gpipe_attn_grad.backward(tensors, grad_output)
196	H	9216	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	self.packages/torch_gpipe_attn_grad.backward(tensors, grad_output)
197	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	self.packages/torch_gpipe_attn_grad.backward(tensors, grad_output)
198	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	self.packages/torch_gpipe_attn_grad.backward(tensors, grad_output)
199	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))
200	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))
201	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))
202	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	attn = (q @ k.transpose(-2, -1))* self.scale
203	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	attn = (q @ k.transpose(-2, -1))* self.scale
204	G	65536	~/mew_mem_1/gpt2_gpipe_memory_profile.py	98	causal_mask = torch.triu(torch.ones(T, T, device=x.device, dtype=torch.bool), diagonal=1)
205	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	100	attn = torch.softmax(attn, dim=-1)
206	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	attn = self.attn_dropout(attn)
207	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	attn = self.attn_dropout(attn)
208	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	out = (attn @ v).transpose(1, 2).contiguous().view(B, T, C)
209	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	out = (attn @ v).transpose(1, 2).contiguous().view(B, T, C)
210	G	393216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	104	out = self.resid_dropout(self.c_proj(out))
211	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))
212	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
213	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
214	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
215	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	119	x = self.c_fc(x)
216	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	120	x = self.act(x)
217	G	393216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	122	x = self.dropout(x)
218	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
219	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))
220	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))
221	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))

Idx	Cat	Size(B)	File	Line	Code
222	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	attn = (q @ k.transpose(-2, -1))* self.scale
223	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	attn = (q @ k.transpose(-2, -1))* self.scale
224	G	65536	~/mew_mem_1/gpt2_gpipe_memory_profile.py	98	causal_mask = torch.triu(torch.ones(T, T, device=x.device, dtype=torch.bool), diagonal=1)
225	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	100	attn = torch.softmax(attn, dim=-1)
226	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	attn = self.attn_dropout(attn)
227	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	attn = self.attn_dropout(attn)
228	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	out = (attn @ v).transpose(1, 2).contiguous().view(B, T, C)
229	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	out = (attn @ v).transpose(1, 2).contiguous().view(B, T, C)
230	G	393216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	104	out = self.resid_dropout(self.c_proj(out))
231	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))
232	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
233	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
234	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
235	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	119	x = self.c_fc(x)
236	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	120	x = self.act(x)
237	G	393216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	122	x = self.dropout(x)
238	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
239	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))
240	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))
241	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))
242	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	attn = (q @ k.transpose(-2, -1))* self.scale
243	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	attn = (q @ k.transpose(-2, -1))* self.scale
244	G	65536	~/mew_mem_1/gpt2_gpipe_memory_profile.py	98	causal_mask = torch.triu(torch.ones(T, T, device=x.device, dtype=torch.bool), diagonal=1)
245	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	100	attn = torch.softmax(attn, dim=-1)
246	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	attn = self.attn_dropout(attn)
247	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	attn = self.attn_dropout(attn)
248	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	out = (attn @ v).transpose(1, 2).contiguous().view(B, T, C)
249	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	out = (attn @ v).transpose(1, 2).contiguous().view(B, T, C)
250	G	393216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	104	out = self.resid_dropout(self.c_proj(out))
251	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))
252	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
253	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
254	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
255	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	119	x = self.c_fc(x)
256	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	120	x = self.act(x)
257	G	393216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	122	x = self.dropout(x)
258	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
259	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))
260	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))
261	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))
262	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	attn = (q @ k.transpose(-2, -1))* self.scale
263	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	attn = (q @ k.transpose(-2, -1))* self.scale

Idx	Cat	Size(B)	File	Line	Code
264	G	65536	~/mew_mem_1/gpt2_gpipe_memory_profile.py	98	causal_mask = torch.triu(torch.ones(T, T, device=x.device, dtype=torch.bool), diagonal=1)
265	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	100	attn = torch.softmax(attn, dim=-1)
266	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	attn = self.attn_dropout(attn)
267	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	attn = self.attn_dropout(attn)
268	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	out = (attn @ v).transpose(1, 2).contiguous().view(B, T, C)
269	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	out = (attn @ v).transpose(1, 2).contiguous().view(B, T, C)
270	G	393216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	104	out = self.resid_dropout(self.c_proj(out))
271	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	x = x + self.attn(self.ln_1(x))
272	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
273	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
274	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
275	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	119	x = self.c_fc(x)
276	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	120	x = self.act(x)
277	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.mlp(self.ln_2(x))
278	H	1572864	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/attn_checkpoint_grad(tensors, grad_output)	260	
279	H	6291456	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/attn_checkpoint_grad(tensors, grad_output)	260	
280	H	9437184	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/attn_checkpoint_grad(tensors, grad_output)	260	
281	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/attn_checkpoint_grad(tensors, grad_output)	260	