

Peak Allocation Events

- A Model Parameters
- B Optimizer States (AdamW)
- C cuBLAS Workspace
- D Pipeline Boundary Copies
- E Gradient Boundary Copies
- F Parameter Gradients
- G Recomputed Activations
- H Final Backward Tensors
- I OTHERS

Idx	Cat	Size(B)	File	Line	Code
1	G	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
2	G	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
3	G	7077888	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
4	G	9216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
5	G	2359296	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
6	G	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
7	G	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
8	G	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
9	G	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
10	G	12288	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
11	G	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
12	G	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
13	G	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
14	G	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
15	G	7077888	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
16	G	9216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
17	G	2359296	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
18	G	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
19	G	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
20	G	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
21	G	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
22	G	12288	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches
23	G	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	193	# checkpoint='always' -> GPipe's built-in activation checkpointing on all micro-batches

Idx	Cat	Size(B)	File	Line	Code
115	B	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
116	B	12288	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
117	B	12288	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
118	B	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
119	B	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
120	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
121	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
122	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
123	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
124	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
125	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
126	B	7077888	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
127	B	7077888	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
128	B	9216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
129	B	9216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
130	B	2359296	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
131	B	2359296	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
132	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
133	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
134	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
135	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
136	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
137	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
138	B	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
139	B	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
140	B	12288	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
141	B	12288	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
142	B	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
143	B	9437184	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
144	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
145	B	3072	~/mew_mem_1/gpt2_gpipe_memory_profile.py	249	loss.backward()
146	D	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	236	# Forward pass through GPipe pipeline
147	C	8519680	~/mew_mem_1/gpt2_gpipe_memory_profile.py	90	B, T, C = x.size()
148	D	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	236	# Forward pass through GPipe pipeline
149	D	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	236	# Forward pass through GPipe pipeline
150	H	9437184	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/this_gpipe_backward(tensors, grad_output)	260	
151	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/this_gpipe_backward(tensors, grad_output)	260	
152	H	9437184	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/this_gpipe_backward(tensors, grad_output)	260	
153	H	12288	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/this_gpipe_backward(tensors, grad_output)	260	
154	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/this_gpipe_backward(tensors, grad_output)	260	
155	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/this_gpipe_backward(tensors, grad_output)	260	
156	H	2359296	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/this_gpipe_backward(tensors, grad_output)	260	
157	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/this_gpipe_backward(tensors, grad_output)	260	
158	E	1572864	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/copy_get_device(prev_stream))	566	
159	H	7077888	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/this_gpipe_backward(tensors, grad_output)	260	

Idx	Cat	Size(B)	File	Line	Code
160	H	9216	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2608	def backward(tensors,
161	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2609	grad_output)
162	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2608	def backward(tensors,
163	H	9437184	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2609	grad_output)
164	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2608	def backward(tensors,
165	H	9437184	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2609	grad_output)
166	H	12288	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2608	def backward(tensors,
167	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2609	grad_output)
168	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2608	def backward(tensors,
169	H	2359296	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2609	grad_output)
170	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2608	def backward(tensors,
171	H	7077888	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2609	grad_output)
172	H	9216	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2608	def backward(tensors,
173	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2609	grad_output)
174	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2608	def backward(tensors,
175	H	9437184	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2609	grad_output)
176	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2608	def backward(tensors,
177	H	9437184	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2609	grad_output)
178	H	12288	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2608	def backward(tensors,
179	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2609	grad_output)
180	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2608	def backward(tensors,
181	H	2359296	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2609	grad_output)
182	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2608	def backward(tensors,
183	H	7077888	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2609	grad_output)
184	H	9216	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2608	def backward(tensors,
185	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2609	grad_output)
186	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2608	def backward(tensors,
187	H	9437184	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2609	grad_output)
188	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gluon_utils.py	2608	def backward(tensors,

Idx	Cat	Size(B)	File	Line	Code
189	H	9437184	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	-packages/torchgpipe/transformer.py:100 def forward(self, x):
190	H	12288	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	-packages/torchgpipe/transformer.py:100 def forward(self, x):
191	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	-packages/torchgpipe/transformer.py:100 def forward(self, x):
192	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	-packages/torchgpipe/transformer.py:100 def forward(self, x):
193	H	2359296	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	-packages/torchgpipe/transformer.py:100 def forward(self, x):
194	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	-packages/torchgpipe/transformer.py:100 def forward(self, x):
195	H	7077888	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	-packages/torchgpipe/transformer.py:100 def forward(self, x):
196	H	9216	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	-packages/torchgpipe/transformer.py:100 def forward(self, x):
197	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	-packages/torchgpipe/transformer.py:100 def forward(self, x):
198	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	-packages/torchgpipe/transformer.py:100 def forward(self, x):
199	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
200	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
201	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
202	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	# Causal mask
203	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	# Causal mask
204	G	65536	~/mew_mem_1/gpt2_gpipe_memory_profile.py	98	attn = (q @ k.transpose(-2, -1))*
205	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	100	self.scale
206	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	attn = attn.masked_fill(causal_mask.unsqueeze
207	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	float('-inf'))
208	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	attn = torch.softmax(attn, dim=-1)
209	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	(source not available)
210	G	393216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	104	out = (attn @ v).transpose(1,
211	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
212	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
213	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
214	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
215	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	119	def forward(self, x):
216	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	120	x = self.c_fc(x)
217	G	393216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	122	x = self.c_proj(x)
218	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
219	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
220	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
221	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
222	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	# Causal mask
223	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	# Causal mask
224	G	65536	~/mew_mem_1/gpt2_gpipe_memory_profile.py	98	attn = (q @ k.transpose(-2, -1))*
225	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	100	self.scale
226	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	attn = attn.masked_fill(causal_mask.unsqueeze
227	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	float('-inf'))
228	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	attn = torch.softmax(attn, dim=-1)
229	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	(source not available)
230	G	393216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	104	out = (attn @ v).transpose(1,
					2).contiguous().view(B, T, C)

Idx	Cat	Size(B)	File	Line	Code
231	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
232	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
233	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
234	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
235	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	119	def forward(self, x):
236	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	120	x = self.c_fc(x)
237	G	393216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	122	x = self.c_proj(x)
238	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
239	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
240	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
241	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
242	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	# Causal mask
243	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	# Causal mask
244	G	65536	~/mew_mem_1/gpt2_gpipe_memory_profile.py	98	attn = (q @ k.transpose(-2, -1))*
					self.scale
245	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	100	attn = attn.masked_fill(causal_mask.unsqueeze
					float('-inf'))
246	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	attn = torch.softmax(attn, dim=-1)
247	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	attn = torch.softmax(attn, dim=-1)
248	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	(source not available)
249	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	(source not available)
250	G	393216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	104	out = (attn @ v).transpose(1,
					2).contiguous().view(B, T, C)
251	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
252	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
253	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
254	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
255	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	119	def forward(self, x):
256	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	120	x = self.c_fc(x)
257	G	393216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	122	x = self.c_proj(x)
258	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
259	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
260	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
261	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
262	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	# Causal mask
263	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	97	# Causal mask
264	G	65536	~/mew_mem_1/gpt2_gpipe_memory_profile.py	98	attn = (q @ k.transpose(-2, -1))*
					self.scale
265	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	100	attn = attn.masked_fill(causal_mask.unsqueeze
					float('-inf'))
266	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	attn = torch.softmax(attn, dim=-1)
267	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	101	attn = torch.softmax(attn, dim=-1)
268	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	(source not available)
269	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	103	(source not available)
270	G	393216	~/mew_mem_1/gpt2_gpipe_memory_profile.py	104	out = (attn @ v).transpose(1,
					2).contiguous().view(B, T, C)
271	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	70	def forward(self, x):
272	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
273	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
274	G	2048	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
275	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	119	def forward(self, x):
276	G	6291456	~/mew_mem_1/gpt2_gpipe_memory_profile.py	120	x = self.c_fc(x)
277	G	1572864	~/mew_mem_1/gpt2_gpipe_memory_profile.py	71	x = x + self.attn(self.ln_1(x))
278	H	1572864	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/ckpt/backward	260	(tensors,
					grad_output)
279	H	6291456	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/site-packages/torch/gpipe/ckpt/backward	260	(tensors,
					grad_output)

Idx	Cat	Size(B)	File	Line	Code
280	H	9437184	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	torch.nn.functional.grad_output)
281	H	3072	~/anaconda3/envs/gpt2_gpipe_mem/lib/python3.10/260	260	torch.nn.functional.grad_output)